

PRAKASH TEKI

Data Engineer | Microsoft Fabric | Dimensional Modeling | Power BI Semantic Models | SQL | DP-600 Certified
Denton, TX | [\(469\) 739-0718](tel:(469)739-0718) | tekiprakash1@gmail.com | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

PROFESSIONAL SUMMARY

Data Engineer with 4+ years building dimensional models, lakehouse and warehouse pipelines, and Power BI semantic models across Microsoft Fabric, BigQuery, and Databricks. Defines data grain, facts, dimensions, and composite keys so models hold from source to report, and evaluates platform trade-offs before build. At Verizon, automates BigQuery revenue reconciliation and validated 125 Teradata to BigQuery ETL processes. Embeds AI through GPT-4, RAG, and LangGraph. Microsoft Certified (DP-600, PL-300).

PROFESSIONAL EXPERIENCE

Verizon, Irving, TX

Data Engineer

Jul 2026 - Present

- Automate Revenue Assurance reconciliation on GCP with BigQuery SQL and Python pipelines that match source and downstream records and surface variances for analyst review.
- Engineer scheduled GCP pipelines landing source data in analytics-ready BigQuery tables that feed Python ML modules.
- Validated 125 migrated Teradata to BigQuery ETL processes with a Python wrapper that runs scripts and reconciles row counts, catching a BETWEEN date-range defect before cutover.
- Built a LookML semantic layer on Oracle with explores and filtered measures, retiring a manual email-based report.

Cates Control Systems, Plano, TX

Data Analyst

May 2026 - Jul 2026

- Modeled Epicor PO header, release, and vendor tables into a 12-table Power BI semantic model with 8 DAX measures.
- Set the fact grain at PO release and built composite keys across two Epicor company codes, so PO aging, open quantity, and past-due measures aggregate correctly across both plant sites.
- Profiled a 50.97 MB Epicor dataset and evaluated Power BI deployment and licensing options, recommending a ~\$210/month Pro plan to scale access to 10 to 15 users.
- Rebuilt an Excel aged-inventory model in Power BI, validating stock segments to within 0.3% of the analyst source.

Holiday Channel, Remote

Data Engineer

Feb 2026 - May 2026

- Built PySpark and SQL medallion pipelines (Bronze, Silver, Gold) that ingest and standardize multi-source content and engagement data, replacing manual analyst prep with daily refreshes.
- Automated a weekly Sheets to PostgreSQL sync in Python with dry-run mode, cutting a 2 to 4 hour job to 2 minutes.
- Delivered Power BI semantic models with DAX measures and row-level security for content and revenue reporting.
- Integrated the GPT-4 API into the content pipeline to auto-generate metadata tags and summaries for each record.

Wisdom Index Advisors, Remote

Data Analyst

Aug 2025 - Dec 2025

- Co-led PostgreSQL cloud database design, defining data requirements and business use cases for client tracking.
- Built a React, Flask, and PostgreSQL app consolidating 37+ financial KPIs into one source, cutting review time 80%.
- Developed SQL stored procedures and Python refresh scripts for client lifetime value, cutting lookup time by 70%.
- Embedded a GPT-4 review agent that flags anomalies in client financials and recommends corrective actions.

University of North Texas, Denton, TX

Data Analyst

Feb 2024 - Aug 2025

- Designed star schema Power BI semantic models with fact and dimension tables and DAX measures tracking transplant operations, regional trends, and capacity across 4,128+ facilities and 3M+ records.
- Automated ingestion and cleansing in Python, SQL, Power Query, and Power Automate with built-in quality checks.
- Built a logistic regression model in scikit-learn and Azure ML to identify kidney transplant success factors.
- Partnered with administrators, IT, and faculty to define KPIs and unify metric definitions across reporting.

Accenture, Hyderabad, India

Data Analyst

Nov 2021 - Dec 2023

- Developed financial dashboards and semantic data models across 5+ business categories, cutting manual reporting 33%.
- Optimized SQL integrating 30+ legacy databases into centralized models, improving retrieval performance by 20%.
- Built ETL validation scenarios with data profiling across OLTP and OLAP sources in enterprise data lakes.
- Partnered with data engineering and architecture teams to migrate and standardize data into governed BI.

TECHNICAL SKILLS

Microsoft Fabric: Lakehouse, Warehouse, OneLake, Notebooks, Data Pipelines, semantic models, medallion architecture (Bronze, Silver, Gold)
Data Modeling: Dimensional modeling, star schema, fact and dimension tables, data grain, composite keys, denormalized views, relational modeling
Power BI: Semantic models, DAX, row-level security, Power Query, Power Automate, KPI dashboards; Looker and LookML; Tableau
SQL and Warehousing: SQL (CTEs, window functions, stored procedures), BigQuery, PostgreSQL, SQL Server, Oracle, Teradata, Snowflake
Data Engineering: Python (Pandas, NumPy), PySpark, Databricks, Hive, HDFS, GCP, Azure Data Factory, ETL/ELT, data quality validation, reconciliation
AI and Machine Learning: GPT-4, LLMs, RAG, LangGraph, prompt engineering, scikit-learn, Azure ML; Git/GitHub, Jira, Agile

PROJECTS

Real-Time Earthquake ELT Pipeline | *Microsoft Fabric, Lakehouse, OneLake, PySpark, Power BI*
• Landed USGS API JSON in a Fabric OneLake Lakehouse through Bronze, Silver, and Gold PySpark notebooks.
• Standardized schema, deduplicated, removed nulls, and type-cast magnitude, depth, and timestamps in the Silver layer.
• Modeled Gold dimensional tables for geospatial hotspot and magnitude trend analysis, surfaced in Power BI dashboards.
• Wrote modular PySpark notebooks with error handling and data quality validation at each layer of the Lakehouse.

E-Commerce ETL Pipeline and BI Model | *Databricks, Apache Spark, SQL, Dimensional Modeling, Power BI*
• Designed 6 modular medallion notebooks (3 dimension, 3 fact) producing Gold fact and dimension tables for BI.
• Engineered Silver business rules for discounts, sale amounts, and currency conversion via a rate dimension join.
• Separated dimension and fact logic into independent notebooks, giving Power BI clean fact-to-dimension joins.
• Published Gold aggregations and denormalized views as Power BI and Tableau sources, plus a Databricks KPI dashboard.

JD E-Commerce Platform Analytics | *Python, Pandas, NumPy, Data Profiling*
• Profiled 549,989 orders, 457,298 users, and 31,868 SKUs with reusable Pandas functions for schema and grain checks.
• Flagged 90.4% and 96.4% missing rates in SKU date fields and documented imputation versus exclusion guidance.

Spotify Listening Patterns Dashboard | *Power BI, DAX, Power Query, Semantic Modeling*
• Modeled 149,860 listening events (2013 to 2024) across 13,839 tracks, 4,113 artists, and 7,948 albums in Power BI.
• Built DAX time-series and YoY measures with platform, shuffle, and skip slicers plus weekday versus weekend views.

IMDb Big Data Warehouse Analytics | *Apache Hive, HiveQL, HDFS, Sqoop*
• Created external Hive tables over HDFS with defined schemas, managing schema evolution under Hive's immutable metadata.
• Wrote HiveQL with window functions, RANK(), and CTEs for genre and director rankings, ingesting RDBMS data via Sqoop.

EDUCATION AND CERTIFICATIONS

M.S. Business Analytics (GPA 3.9/4.0), University of North Texas, Denton, TX	Jan 2024 - Dec 2025
B.Com Honours (GPA 3.6/4.0), Osmania University, Hyderabad, India	Jul 2018 - Aug 2021
Microsoft Certified: Fabric Analytics Engineer Associate (DP-600) Power BI Data Analyst Associate (PL-300)	
DataCamp: Data Warehousing Concepts, Database Design, Introduction to Relational Databases in SQL, Understanding Cloud Computing	